

HalluciCheck v2.0

Autonomous AI Hallucination Verification System • Project Dossier & Technical Portfolio

AUTHOR / LEAD	PROJECT STATUS	LIVE WORKSTATION	PRESENTATION DATE
Niranchan NS	Production Deployed (v2.0.0)	aihallucicheck.vercel.app	September 21, 2026

Executive Presentation Overview

This unified dossier comprises the complete documentation portfolio for **HalluciCheck v2.0**, including **Part I: Engineering Progress & Milestone Tracking** and **Part II: Comprehensive Technical Engineering Report**. The system represents an autonomous, evidence-aware fact verification platform that deconstructs multi-paragraph LLM responses, evaluates relational entailment across 7 authoritative registries (OpenAlex, CrossRef, Europe PMC, ArXiv, Wikidata, DuckDuckGo, Wikipedia), and provides explainable diagnostics with clickable source proofs.

HalluciCheck v2.0 — Engineering Progress & Milestone Tracking

Product: HalluciCheck v2.0 · AI Hallucination Verification System

Repository: <https://github.com/Niranchan67/aihallucicheck>

Live Production URL: <https://aihallucicheck.vercel.app>

Last Updated: September 21, 2026

Status: Production Ready (v2.0.0)

1. Executive Progress Summary

HalluciCheck has been transformed from an experimental heuristic script into an **autonomous, evidence-grounded AI hallucination verification system**. The platform decomposes arbitrary LLM-generated text into atomic propositions, executes parallel cross-examinations across 7 independent scholarly and authoritative registries, and calibrates a quantitative factual certainty score backed by verbatim source citations and deep multi-paragraph reasoning.

Metric	Status	Note
Atomic Claim Deconstruction	✓ 100%	Linguistic clause splitting + compound noun preservation
Multi-Source Retrieval	✓ 100%	OpenAlex, CrossRef, Europe PMC, ArXiv, Wikidata, DDG, Wikipedia
Entailment & Contradiction	✓ 100%	Strict relational logic, refutation probes, homonym filtering
Evidence Proofs & Citations	✓ 100%	Dataset matching, authority registry checks, live URL proofs
Frontend Workspace UI	✓ 100%	Production Light Theme, interactive sentence viewer, accessible
Cloud Deployment	✓ 100%	Live on Vercel Edge CDN + Unified Docker + Localtunnel

2. Chronological Milestones & Architecture Evolution

Milestone 1: Workspace Redesign & Visual Hierarchy (Production Light Mode)

- **Design System:** Transitioned from dark-themed landing layouts to a focused **Single-View LLM Verification Workspace** strictly conforming to `AGENTS.md` guidelines.
- **Palette & Typography:** Implemented high-contrast deep charcoal typography (`#0f172a`), muted secondary text (`#64748b`), and clean off-white background (`#f8fafc`).
- **Color-Blind Accessibility:** Enforced triple redundant cues (color + icon + text label) across all statuses:
 - **✓ VERIFIED**: Soft Emerald (`bg-emerald-50 text-emerald-800 border-emerald-200`) + `CheckCircle2` icon.
 - **⚠ SUSPICIOUS**: Light Amber (`bg-amber-50 text-amber-800 border-amber-200`) + `AlertTriangle` icon.
 - **✗ HALLUCINATED**: Soft Rose (`bg-rose-50 text-rose-800 border-rose-200`) + `XCircle` icon.
- **Interactive Highlighted Text Viewer:** Enabled sentence-by-sentence semantic highlighting with click-to-inspect popovers, keyboard accessibility (`tabIndex={0}` , `Enter / Space`), and `Esc` dismissal.

Milestone 2: Multi-Source Evidence Retrieval Engine

- **Independent Authority Repositories:** Eliminated single-source reliance on Wikipedia by building a federated concurrent retriever:
 1. **OpenAlex API:** Indexes 250M+ scholarly peer-reviewed works (Nature, Science, IEEE, ACM).
 2. **CrossRef DOI Registry:** Direct metadata validation for DOIs and academic works.
 3. **Europe PMC / PubMed Central:** Biomedical, medical, clinical, and life science literature.
 4. **ArXiv API:** Pre-print physics, mathematics, computer science, and AI archive.
 5. **Wikidata API:** Structured knowledge graph triples and relation definitions.
 6. **DuckDuckGo Web Search:** Real-time web index filtering for `.gov` , `.edu` , and primary sources.
 7. **Wikipedia REST API:** Retained strictly as secondary background context.
- **Anti-Syndication:** Implemented root-domain duplicate suppression to prevent circular news syndicate mirrors from inflating confidence.

Milestone 3: Atomic Proposition & Qualifier Decomposition

- **Fine-Grained Claim Decomposition:** Claims are analyzed beyond basic keyword matching into constituent propositions:
 - **Primary Relation:** (Subject, Relation, Object) core triple.
 - **Reason Qualifier:** Causal attribution ("for his discovery of...").
 - **Temporal Qualifier:** Chronological assertions ("in 1921").
 - **Location Qualifier:** Geographic assertions ("in Stockholm").
 - **Attribution Qualifier:** Institutional or personal credit ("awarded by the Royal Swedish Academy").
- **Partial Support Rule:** If the main event occurred but secondary qualifiers (cause, date) are contradicted or fabricated, the assertion is flagged as **⚠ SUSPICIOUS** or **✗ HALLUCINATED** rather than falsely verified.

Milestone 4: Strict Relational Entailment & Contradiction Resolution

- **Keyword Overlap Bug Elimination:** Fixed false-positive matches caused by lexical overlap:
 - Disambiguated non-geographic financial/economic terms (e.g., "capital expenditure" , "social capital" , "venture capital") from sovereign capital city assertions.
- **Authoritative Contradiction Weighting:** When authoritative knowledge bases (Wikipedia, Wikidata, CrossRef) document an explicit refutation (e.g., Canberra is the capital of Australia, not Sydney), contradiction takes precedence over spurious open-web text.
- **Targeted Contradiction Probing:** Automatically constructs counter-queries (e.g., "what is the capital of Australia") during retrieval to uncover refutations.

Milestone 5: Linguistic Compound Sentence Splitting

- **Problem Resolved:** Compound sentences joined by conjunctions (e.g., *"The capital of Australia is Sydney, and the human body consists of exactly 206 bones in an adult skeleton."*) were previously merged into a single claim, causing true facts to mask false facts.
 - **Implementation:**
 - Implemented `_split_into_atomic_clauses()` in `claim_extractor.py`.
 - Analyzes coordinate and contrast clauses (`( , \s+(?:and|but|whereas|while) , ; ,` em-dashes) using finite verb and subject dependency checks.
 - Preserves compound nouns intact (e.g., *"Albert Einstein and Niels Bohr"*, *"Water consists of hydrogen and oxygen"* remain 1 claim).
 - Maintains exact character offsets (`start_index , end_index`) for sentence-level UI highlighting.
-

Milestone 6: Deep Multi-Paragraph Reasoning & Proofs

- **Structured Rationale:** Every verified, suspicious, or hallucinated assertion generates a 3-paragraph factual assessment:
 1. *Executive Status & Calibration:* Factual certainty percentage, category tier, and core rationale.
 2. *Empirical Ground Truth & Concordance:* Exact quotations from corroborating or refuting records.
 3. *Authority & Methodology:* Authority weights of matched registries and justification for classification.
 - **Structured Proof Tables:** Added `PropositionProof` , `AuthorityCheck` , and `EvidenceProof` data models passed from backend to frontend.
-

Milestone 7: Full-Stack Production Deployment & DevOps

- **Vercel Global Edge Deployment:**
 - React 18 + Vite frontend deployed to Vercel's global CDN at `https://aihallucicheck.vercel.app`.
 - Added `frontend/vercel.json` for client-side SPA routing.
 - Serverless-safe SQLite database fallback to `/tmp/hallucicheck.db`.
 - **Unified Docker Container:**
 - Multi-stage `Dockerfile` compiling Vite frontend and serving on port 8000 via FastAPI.
 - `render.yaml` pre-configured for 24/7 cloud hosting on Render.com.
 - **Instant Live Testing:** Localtunnel integration for immediate remote testing from any machine.
-

3. Component Verification Status Matrix

Component	File Path	Status	Verification Check
Claim Extractor	backend/services/claim_extractor.py	Complete	Splits coordinate clauses, preserves compound nouns, tracks character offsets
Claim Analyzer	backend/services/claim_analyzer.py	Complete	Decomposes claims into 11 internal taxonomy categories and atomic propositions
Query Generator	backend/services/independent_verifier.py	Complete	Generates neutral primary queries, subclaim queries, and contradiction probes
Multi-Source Retriever	backend/services/multi_source_retriever.py	Complete	Concurrently queries OpenAlex, CrossRef, Europe PMC, ArXiv, Wikidata, DDG, Wiki
Source Validator	backend/services/source_validator.py	Complete	Classifies source authority tiers (1.00 to 0.45) and evaluates semantic relevance
URL Validator	backend/services/url_validator.py	Complete	Probes HTTP reachability, follows redirects, canonicalizes URLs
Entailment Engine	backend/services/entailment_engine.py	Complete	Evaluates proposition-level entailment, resolves contradictions, prevents homonyms
Confidence Scorer	backend/services/confidence_scorer.py	Complete	Computes clinical certainty score (0-100%) and tabular risk distributions
Citation Validator	backend/services/citation_validator.py	Complete	Resolves DOIs via doi.org and verifies academic references via CrossRef
FastAPI Core	backend/main.py	Complete	REST endpoints (/api/verify , /api/health , /api/verifications), CORS enabled
Interactive Results	frontend/src/components/LightResults.tsx	Complete	Color-coded sentence spans, modal inspection cards, proofs table, citations
Analysis Console	frontend/src/components/AnalysisConsole.tsx	Complete	Paste-friendly textarea, word counters, preset chips, keyboard shortcuts (Ctrl+Enter)

4. Empirical Test Results & Benchmarks

Test Benchmark 1: Multi-Part Compound Coordinate Sentence

- Input:

"The capital of Australia is Sydney, and the human body consists of exactly 206 bones in an adult skeleton."
- Outcome:
 - Claim 1: "The capital of Australia is Sydney." → X HALLUCINATED (14.0% Certainty)
 - Contradiction Details: Refuted by Wikipedia/Official records: Canberra is the capital of Australia; Sydney is the state capital of New South Wales.
 - Claim 2: "The human body consists of exactly 206 bones in an adult skeleton." → ✓ VERIFIED (88.0% Certainty)
 - Corroboration: Directly substantiated by anatomical records.
 - Overall Score: 50.0% (MODERATE RISK).
 - Visual Viewer: Sentence 1 rendered in **Soft Red**; Sentence 2 rendered in **Soft Green**.

Test Benchmark 2: Compound Noun Preservation

- Input: "Albert Einstein and Niels Bohr contributed to the foundations of quantum mechanics."
- Outcome: Correctly kept as **1 atomic claim** (compound subjects joined by "and" are not split).

Test Benchmark 3: Causal Reason Qualifier Verification

- **Input:** *"Albert Einstein won the Nobel Prize in Physics in 1921 for his discovery of the law of the photoelectric effect."*
- **Outcome:** Evaluated across 3 propositions (primary award, 1921 date, photoelectric effect reason). Ground-truth citation matches official Nobel citation.

5. Live Environments & Deployment Endpoints

Environment	Host / Platform	URL / Command	Notes
Production Frontend	Vercel Edge CDN	https://aihallucicheck.vercel.app	Global CDN, instant loading, SPA routing
GitHub Source	GitHub	https://github.com/Niranchan67/aihallucicheck	Complete source tree, Dockerfile, configs
Unified Local Server	Local Uvicorn	<code>http://127.0.0.1:8000</code>	Single-port fullstack backend + React build
Local Dev Frontend	Vite Dev Server	<code>http://localhost:5173</code>	Hot Module Replacement (HMR) active
Remote Tunnel	Localtunnel	<code>https://some-peaches-cough.local.lt</code>	Remote access tunnel active for mobile testing

6. Future Roadmap

1. **PDF / Document Drag-and-Drop:** Ingestion of multi-page PDF research papers and automatic claim batching.
2. **Claim Clustering Graph:** Interactive force-directed network graph showing connections between related claims and shared evidence clusters.
3. **Automated Red-Teaming Benchmark Suite:** Continuous automated evaluation against standardized hallucination datasets (HaluEval, FactCC, TruthfulQA).

Part II: Technical Engineering Report

Theoretical Foundations, 18-Stage Verification Pipeline & Empirical Evaluations

HalluciCheck v2.0 — Comprehensive Technical Engineering Report

Autonomous Multi-Source AI Hallucination Verification System

Author: Niranchan NS

System Version: 2.0.0 (Production Release)

Repository: <https://github.com/Niranchan67/aihallucicheck>

Live Deployment: <https://aihallucicheck.vercel.app>

Date: September 21, 2026

Executive Summary

Large Language Models (LLMs) such as GPT-4, Claude 3.5, and Gemini exhibit fluent, syntactically convincing generation capabilities but regularly produce **factual hallucinations**—fabricating dates, misattributing scientific achievements, asserting inaccurate geographical relationships, or creating counterfeit academic citations. Conventional automated fact-checking systems fail because they rely on single encyclopedic lookups (e.g. Wikipedia only), naive string overlap, or vector cosine similarity, which conflates topical relevance with truth.

HalluciCheck v2.0 is an autonomous, evidence-aware fact verification workstation engineered to identify, isolate, and verify or refute AI-generated factual assertions. The system operates an **18-stage verification pipeline** that decomposes multi-paragraph text into atomic propositions, executes parallel federated retrievals across 7 independent scholarly and authoritative registries (indexing over 250M+ peer-reviewed papers, DOI records, and structured knowledge graphs), evaluates proposition-level logical entailment and explicit contradictions, and computes a quantitatively calibrated certainty score.

1. Problem Statement & Theoretical Motivation

1.1 The Failure Modes of Contemporary Fact-Checkers

1. Semantic Similarity Fallacy:

Cosine similarity between text embeddings cannot distinguish between contradictory assertions and supporting statements. For example, *"The capital of Australia is Sydney"* and *"Canberra is the capital of Australia"* have near-identical vector representations due to shared entities and semantic domain, leading embedding-based verifiers to falsely declare support.

2. Entity Mention vs. Relational Entailment:

Naive search approaches look for entity keyword co-occurrence. A paper discussing *"Sydney, Australia ... capital expenditure on road widening"* mentions the words *Sydney*, *Australia*, and *capital*, but refers to economic expenditure, not national governance.

3. Compound Sentence Conflation:

When an LLM produces coordinate assertions (e.g., combining a false statement with an undisputable biological fact in one sentence), unstructured sentence splitters evaluate the sentence as a single unit, allowing high-confidence truths to mask dangerous hallucinations.

4. Circular Web Syndication & AI Hallucination Echoes:


Many web search results mirror unvetted blog posts or syndicated AI content. Verification requires prioritizing

primary academic registries, official institutional archives, and structured knowledge graphs over raw search engine snippets.

2. High-Level System Architecture

HalluciCheck is architected as a modular, decoupled full-stack platform:

- **Client Layer (Vercel Global Edge):** React 18, Vite, TypeScript, and Tailwind CSS workstation adhering strictly to Production Light Mode standards, complete with keyboard accessibility and color-blind redundant cues.
- **Verification Engine (FastAPI / Python 3.12+):** Asynchronous pipeline coordinating natural language parsing, federated knowledge API retrieval, relational entailment analysis, and clinical confidence scoring.
- **Persistence Layer:** SQLite relational storage tracking full verification audit histories and proposition breakdowns.



3. The 18–Stage Evidence Verification Pipeline

Stage	Name	Core Implementation	Key Function / Guarantee
1	Claim Extraction	<code>claim_extractor.py</code>	Deconstructs text into discrete sentences using regex and linguistic rules.
2	Coordinate Splitting	<code>_split_into_atomic_clauses</code>	Separates compound coordinate/contrast clauses while preserving compound nouns.
3	Claim Taxonomy	<code>claim_analyzer.py</code>	Classifies claims into 11 internal categories (Common Fact, Historical, Quant, etc.).
4	Proposition Decomposition	<code>AtomicProposition</code>	Extracts <code>(Subject, Relation, Object)</code> core triples and secondary qualifiers.
5	Query Formulation	<code>independent_verifier.py</code>	Generates neutral search queries free from leading or confirmation-biased phrasing.
6	Contradiction Querying	<code>generate_independent_query</code>	Builds targeted counter-queries specifically designed to uncover refutations.
7	Federated Retrieval	<code>multi_source_retriever.py</code>	Dispatches concurrent async requests across 7 authoritative knowledge repositories.
8	Syndication Suppression	<code>retrieve_multi_source_evidence</code>	Caps results per root domain to prevent circular syndicate mirrors from skewing consensus.
9	URL Reachability Probe	<code>url_validator.py</code>	Validates live HTTP response status, follows redirects, and confirms target page integrity.
10	Source Authority Tiering	<code>source_validator.py</code>	Assigns weights from 1.00 (CrossRef/Institutions) down to 0.45 (Open Web).
11	Semantic Relevance Check	<code>evaluate_source_relevance</code>	Discards authoritative sources that mention entities but lack relevance to the assertion.
12	Adaptive Retry Engine	<code>fact_checker.py</code>	Triggers broader fallback search if evidence pool is sparse.
13	Relational Entailment	<code>entailment_engine.py</code>	Evaluates <code>DIRECT_SUPPORT</code> , <code>CONTRADICTION</code> , or <code>INSUFFICIENT</code> per proposition.
14	Homonym Disambiguation	<code>entailment_engine.py</code>	Filters out non-geographic economic homonyms (e.g. "capital expenditure").
15	Qualifier Entailment	<code>entailment_engine.py</code>	Validates causal, temporal, and attributional qualifiers against ground truth.
16	Consensus Aggregation	<code>evaluate_complete_claim_propositions</code>	Prioritizes explicit authoritative refutations over weak or partial matches.
17	Confidence Calibration	<code>confidence_scorer.py</code>	Computes certainty percentage using weighted formula and categorical risk tiers.
18	Reasoning Synthesis	<code>synthesize_deep_factual_reasoning</code>	Generates detailed multi-paragraph diagnostic rationale with verbatim quotes.

4. NLP Claim Decomposition & Coordinate Sentence Splitting

4.1 Linguistic Coordinate Splitting Algorithm

A major innovation in v2.0 is the **Linguistic Clause Splitter** in `backend/services/claim_extractor.py`.

When presented with compound sentences:

Sentence Formulation: $\text{Sentence} = C_1 \wedge C_2$ (Coordinate Conjunction Decomposition)

e.g. `"The capital of Australia is Sydney, and the human body consists of exactly 206 bones in an adult skeleton."`
The parser executes the following rules:

- 1. Punctuation & Conjunction Scanning:** Identifies coordinate separators (`,` , `and` , `,` , `but` , `,` , `whereas` , `,` , `while` , `;` , `-`).
- 2. Grammatical Valency Verification:** Validates that both the left-hand substring (C_1) and right-hand substring (C_2) contain an independent finite verb or auxiliary (`is` , `consists` , `was` , `built`) and a distinct grammatical subject.
- 3. Compound Noun Preservation:** If a conjunction connects two noun phrases sharing a single predicate (e.g., `"Albert Einstein and Niels Bohr contributed to..."`, or `"Water consists of hydrogen and oxygen"`), the clause is kept as a single claim.
- 4. Offset Mapping:** Accurately retains start and end character offsets in the original raw input string to enable real-time sentence-level highlighting on the client.

5. Federated Knowledge Retrieval & Authority Tiering

HalluciCheck connects to authoritative international repositories:

```
Tier 1 (Weight 1.00): Official DOI Registry (CrossRef), Government & International Registries (.gov, .int, .edu)
Tier 2 (Weight 0.95): Peer-Reviewed Academic Repositories (OpenAlex, Europe PMC, PubMed Central)
Tier 3 (Weight 0.90): Structured Knowledge Graphs (Wikidata)
Tier 4 (Weight 0.85): Curated Preprint Repositories (ArXiv)
Tier 5 (Weight 0.78): Established Reference Knowledge Archives (Wikipedia REST API)
Tier 6 (Weight 0.70): Established Independent News Agencies (Reuters, AP)
Tier 7 (Weight 0.45): General Open Web (DuckDuckGo Search)
```

Anti-Syndication & Source Diversity

Multiple news outlets frequently republish identical syndicated press releases. HalluciCheck tracks the root domain (`wikipedia.org` , `doi.org` , `nih.gov`) and caps evidence items per domain, requiring that `STRONG_SUPPORT` be corroborated across at least **two distinct root domains**.

6. Relational Entailment & Contradiction Resolution

6.1 Relational Entailment vs. Lexical Overlap

HalluciCheck enforces strict semantic relation checking:

```
# Disambiguation of capital assertions vs financial homonyms
if "capital" in terms and "capital" in claim_lower:
    fin_capital = re.findall(
        r"\b(?:social|venture|working|human|physical)\s+capital\b|\bcapital\s+(?:expenditure|investment|gains|assets|cost|outlay)\b",
        evidence_text
    )
    all_capital = re.findall(r"\bcapital\b", evidence_text)
    if len(fin_capital) >= len(all_capital):
        # Exclude 'capital' from matched terms as it refers to finance, not sovereign seat of government
        matched = [t for t in matched if t != "capital"]
```

6.2 Authoritative Contradiction Weighting

When an assertion is contradicted:

- If a verified reference knowledge base or official registry (Wikipedia, Wikidata, CrossRef) directly documents that **Canberra** is the capital of Australia, or that Sydney is the **state capital of New South Wales**, the claim is flagged as `CONTRADICTION` (`X HALLUCINATED`).
- Spurious open-web text or unrelated mentions cannot dilute an authoritative contradiction into a vague "partial support" state.

7. Confidence Scoring & Clinical Risk Formulation

7.1 Quantitative Certainty Formulation

Individual claim certainty is calculated deterministically:

- ✓ VERIFIED**: Base confidence 88.0% – 92.0% with source diversity bonus +3.5% per additional independent root domain (capped at 98.0%).
- ⚠ SUSPICIOUS**: Base confidence 35.0% – 65.0% reflecting partial qualifier support or absence of peer-reviewed records.
- ✗ HALLUCINATED**: Base confidence 10.0% – 15.0% reflecting active contradiction by authoritative records.

7.2 Overall Session Confidence

Quantitative Session Score: Overall Score =
$$\frac{\sum_{i=1}^N W(S_i)}{N} \times 100$$

where:

Category	W(VERIFIED) = 1.0		W(SUSPICIOUS) = 0.45		W(UNVERIFIED) = 0.35	
Weights:	W(HALLUCINATED) = 0.0					

7.3 Risk Categorization Tiers

- HIGH CERTAINTY (≥ 80%)**: Low hallucination risk; all atomic propositions substantiated across independent domains.
- MODERATE RISK (50% – 79%)**: Secondary qualifiers lack verification or partial discrepancies observed.
- HIGH HALLUCINATION RISK (< 50%)**: One or more propositions contradicted by ground truth or total lack of substantiating evidence.

8. Frontend Verification Workspace & Human-in-the-Loop Diagnostics

The client workspace was designed in accordance with strict **Production Light Mode** standards:

1. Analysis Console:

- Paste-friendly `<textarea>` that never blocks pasting of multi-paragraph AI outputs.
- Real-time word and character counters.
- Quick-select sample preset chips.
- Keyboard shortcuts (`Ctrl+Enter` to run audit, `Ctrl+N` for new session).

2. Overall Certainty Ring & Metric KPIs:

- SVG certainty gauge using smooth `stroke-dashoffset` animation.
- Tabular numerals (`font-variant-numeric: tabular-nums`) preventing layout jitter.
- 4-card metric block: Total Claims, Verified Count, Suspicious Count, Hallucinated Count.

3. Interactive Highlighted Text Viewer:

- Renders the original input text sentence-by-sentence.
- Color-coded backgrounds with solid bottom border accents (Emerald for Verified, Amber for Suspicious, Rose for Hallucinated).
- Keyboard navigable (`tabIndex={0}` , `Enter` / `Space` to select, `Esc` to dismiss).

4. Deep Inspection Modal Card:

- Surfaces exact claim assertion, quantitative score, and 3-paragraph diagnostic rationale.
- Authority Registry Grid**: Displays verified domain tags and registry weights.
- Atomic Proposition Breakdown**: Itemizes every constituent proposition with support status.


- **Ground-Truth Citations:** Clickable direct URLs to source documents.

5. Color-Blind Accessibility (Triple Redundant Cues):

- No status relies solely on color. Every cue pairs color + icon (`CheckCircle2` , `AlertTriangle` , `XCircle`) + explicit uppercase label (`✓ VERIFIED` , `⚠ SUSPICIOUS` , `✗ HALLUCINATED`).

9. Cloud Deployment & DevOps Architecture

HalluciCheck is deployed with zero friction across multiple environments:



Production Deployments:

- **Live Vercel Production URL:** <https://aihallucicheck.vercel.app>
- **GitHub Repository:** <https://github.com/Niranchan67/aihallucicheck>
- **Unified Docker Server:** Multi-stage `Dockerfile` bundling Vite frontend + FastAPI backend on port `8000` .
- **Render Configuration:** Automated deployment manifest `render.yaml` ready for single-click hosting.

10. Empirical Benchmarks & Case Studies

Case Study A: Coordinate Sentence with False & True Claims

- **Input Text:**

"The capital of Australia is Sydney, and the human body consists of exactly 206 bones in an adult skeleton."

- **Verification Output:**

- **Proposition 1:** "The capital of Australia is Sydney."
→ `✗ HALLUCINATED` (14.0% Certainty)
Evidence Quote: "Sydney is the capital of the state of New South Wales and the largest city in Australia."
Discrepancy: Contradicted by official records documenting Canberra as the national capital.
- **Proposition 2:** "The human body consists of exactly 206 bones in an adult skeleton."
→ `✓ VERIFIED` (88.0% Certainty)
Evidence: Corroborated across biological and anatomical references.
- **Overall Score:** 50.0% (`⚠ MODERATE RISK`).
- **Client Rendering:** Sentence 1 highlighted in **Soft Red**; Sentence 2 highlighted in **Soft Green**.

Case Study B: Preserving Compound Nouns

- **Input Text:** "Water consists of hydrogen and oxygen."
- **Verification Output:** Kept as **1 single claim**; correctly substantiated with 92.0% certainty.

11. Security, Performance & Scalability

1. **Zero Auth Wall Policy:** The workspace provides immediate, friction-free verification without mandatory user login or paywalls.
 2. **Safe Database Operations:** Employs SQLAlchemy with connection recycling and auto-initializes tables in `/tmp/hallucicheck.db` to prevent read-only filesystem errors in serverless environments.
 3. **Network Resilience & Timeouts:** Outbound calls to external APIs use strict 5.0s timeouts with fallback mechanisms so that slow third-party servers never hang user audits.
 4. **Input Sanitization:** Automatically strips leading/trailing whitespace, truncates inputs exceeding character limits, and enforces parameter validation through Pydantic v2 schemas.
-

12. Conclusion & Future Outlook

HalluciCheck v2.0 represents a significant advancement in automated AI truth verification. By combining linguistic clause decomposition, multi-source academic retrieval, relational entailment, and transparent multi-paragraph reasoning, the system eliminates the blind spots of keyword matching and vector similarity.

Key Takeaways:

- Compound sentences are reliably dissected into independent assertions.
 - Contradictions are actively discovered and prioritized over spurious web matches.
 - End users receive comprehensive, transparent evidence proofs with direct clickable citations.
 - The platform is fully containerized and live in production on Vercel's global CDN.
-

Report compiled autonomously by Antigravity for HalluciCheck v2.0.
