

HalluciCheck v2.0

Autonomous Multi-Source AI Hallucination Verification • Comprehensive Technical Report

AUTHOR

Niranchan NS

SYSTEM ARCHITECTURE

18-Stage Verification
Pipeline

LIVE PRODUCTION URL

aihallucicheck.vercel.app

EVALUATION DATE

September 21, 2026

HalluciCheck v2.0 — Comprehensive Technical Engineering Report

Autonomous Multi-Source AI Hallucination Verification System

Author: Niranchan NS

System Version: 2.0.0 (Production Release)

Repository: <https://github.com/Niranchan67/aihallucicheck>

Live Deployment: <https://aihallucicheck.vercel.app>

Date: September 21, 2026

Executive Summary

Large Language Models (LLMs) such as GPT-4, Claude 3.5, and Gemini exhibit fluent, syntactically convincing generation capabilities but regularly produce **factual hallucinations**—fabricating dates, misattributing scientific achievements, asserting inaccurate geographical relationships, or creating counterfeit academic citations. Conventional automated fact-checking systems fail because they rely on single encyclopedic lookups (e.g. Wikipedia only), naive string overlap, or vector cosine similarity, which conflates topical relevance with truth.

HalluciCheck v2.0 is an autonomous, evidence-aware fact verification workstation engineered to identify, isolate, and verify or refute AI-generated factual assertions. The system operates an **18-stage verification pipeline** that decomposes multi-paragraph text into atomic propositions, executes parallel federated retrievals across 7 independent scholarly and authoritative registries (indexing over 250M+ peer-reviewed papers, DOI records, and structured knowledge graphs), evaluates proposition-level logical entailment and explicit contradictions, and computes a quantitatively calibrated certainty score.

1. Problem Statement & Theoretical Motivation

1.1 The Failure Modes of Contemporary Fact-Checkers

1. Semantic Similarity Fallacy:

Cosine similarity between text embeddings cannot distinguish between contradictory assertions and supporting statements. For example, *"The capital of Australia is Sydney"* and *"Canberra is the capital of Australia"* have near-identical vector representations due to shared entities and semantic domain, leading embedding-based verifiers to falsely declare support.

2. Entity Mention vs. Relational Entailment:

Naive search approaches look for entity keyword co-occurrence. A paper discussing *"Sydney, Australia ... capital expenditure on road widening"* mentions the words *Sydney*, *Australia*, and *capital*, but refers to economic expenditure, not national governance.

3. Compound Sentence Conflation:

When an LLM produces coordinate assertions (e.g., combining a false statement with an undisputable biological fact in one sentence), unstructured sentence splitters evaluate the sentence as a single unit, allowing high-confidence truths to mask dangerous hallucinations.


4. Circular Web Syndication & AI Hallucination Echoes:

Many web search results mirror unvetted blog posts or syndicated AI content. Verification requires prioritizing primary academic registries, official institutional archives, and structured knowledge graphs over raw search engine snippets.

2. High-Level System Architecture

HalluciCheck is architected as a modular, decoupled full-stack platform:

- **Client Layer (Vercel Global Edge):** React 18, Vite, TypeScript, and Tailwind CSS workstation adhering strictly to Production Light Mode standards, complete with keyboard accessibility and color-blind redundant cues.
- **Verification Engine (FastAPI / Python 3.12+):** Asynchronous pipeline coordinating natural language parsing, federated knowledge API retrieval, relational entailment analysis, and clinical confidence scoring.
- **Persistence Layer:** SQLite relational storage tracking full verification audit histories and proposition breakdowns.



3. The 18–Stage Evidence Verification Pipeline

Stage	Name	Core Implementation	Key Function / Guarantee
1	Claim Extraction	<code>claim_extractor.py</code>	Deconstructs text into discrete sentences using regex and linguistic rules.
2	Coordinate Splitting	<code>_split_into_atomic_clauses</code>	Separates compound coordinate/contrast clauses while preserving compound nouns.
3	Claim Taxonomy	<code>claim_analyzer.py</code>	Classifies claims into 11 internal categories (Common Fact, Historical, Quant, etc.).
4	Proposition Decomposition	<code>AtomicProposition</code>	Extracts <code>(Subject, Relation, Object)</code> core triples and secondary qualifiers.
5	Query Formulation	<code>independent_verifier.py</code>	Generates neutral search queries free from leading or confirmation-biased phrasing.
6	Contradiction Querying	<code>generate_independent_query</code>	Builds targeted counter-queries specifically designed to uncover refutations.
7	Federated Retrieval	<code>multi_source_retriever.py</code>	Dispatches concurrent async requests across 7 authoritative knowledge repositories.
8	Syndication Suppression	<code>retrieve_multi_source_evidence</code>	Caps results per root domain to prevent circular syndicate mirrors from skewing consensus.
9	URL Reachability Probe	<code>url_validator.py</code>	Validates live HTTP response status, follows redirects, and confirms target page integrity.
10	Source Authority Tiering	<code>source_validator.py</code>	Assigns weights from 1.00 (CrossRef/Institutions) down to 0.45 (Open Web).
11	Semantic Relevance Check	<code>evaluate_source_relevance</code>	Discards authoritative sources that mention entities but lack relevance to the assertion.
12	Adaptive Retry Engine	<code>fact_checker.py</code>	Triggers broader fallback search if evidence pool is sparse.
13	Relational Entailment	<code>entailment_engine.py</code>	Evaluates <code>DIRECT_SUPPORT</code> , <code>CONTRADICTION</code> , or <code>INSUFFICIENT</code> per proposition.
14	Homonym Disambiguation	<code>entailment_engine.py</code>	Filters out non-geographic economic homonyms (e.g. "capital expenditure").
15	Qualifier Entailment	<code>entailment_engine.py</code>	Validates causal, temporal, and attributional qualifiers against ground truth.
16	Consensus Aggregation	<code>evaluate_complete_claim_propositions</code>	Prioritizes explicit authoritative refutations over weak or partial matches.
17	Confidence Calibration	<code>confidence_scorer.py</code>	Computes certainty percentage using weighted formula and categorical risk tiers.
18	Reasoning Synthesis	<code>synthesize_deep_factual_reasoning</code>	Generates detailed multi-paragraph diagnostic rationale with verbatim quotes.

4. NLP Claim Decomposition & Coordinate Sentence Splitting

4.1 Linguistic Coordinate Splitting Algorithm

A major innovation in v2.0 is the **Linguistic Clause Splitter** in `backend/services/claim_extractor.py`.

When presented with compound sentences:

Sentence Formulation: $\text{Sentence} = C_1 \wedge C_2$ (Coordinate Conjunction Decomposition)

e.g. `"The capital of Australia is Sydney, and the human body consists of exactly 206 bones in an adult skeleton."`
The parser executes the following rules:

1. **Punctuation & Conjunction Scanning:** Identifies coordinate separators (`,` , `and` , `,` , `but` , `,` , `whereas` , `,` , `while` , `;` , `-`).
2. **Grammatical Valency Verification:** Validates that both the left-hand substring (C_1) and right-hand substring (C_2) contain an independent finite verb or auxiliary (`is` , `consists` , `was` , `built`) and a distinct grammatical subject.
3. **Compound Noun Preservation:** If a conjunction connects two noun phrases sharing a single predicate (e.g., *"Albert Einstein and Niels Bohr contributed to..."*, or *"Water consists of hydrogen and oxygen"*), the clause is kept as a single claim.
4. **Offset Mapping:** Accurately retains start and end character offsets in the original raw input string to enable real-time sentence-level highlighting on the client.

5. Federated Knowledge Retrieval & Authority Tiering

HalluciCheck connects to authoritative international repositories:

```
Tier 1 (Weight 1.00): Official DOI Registry (CrossRef), Government & International Registries (.gov, .int, .edu)
Tier 2 (Weight 0.95): Peer-Reviewed Academic Repositories (OpenAlex, Europe PMC, PubMed Central)
Tier 3 (Weight 0.90): Structured Knowledge Graphs (Wikidata)
Tier 4 (Weight 0.85): Curated Preprint Repositories (ArXiv)
Tier 5 (Weight 0.78): Established Reference Knowledge Archives (Wikipedia REST API)
Tier 6 (Weight 0.70): Established Independent News Agencies (Reuters, AP)
Tier 7 (Weight 0.45): General Open Web (DuckDuckGo Search)
```

Anti-Syndication & Source Diversity

Multiple news outlets frequently republish identical syndicated press releases. HalluciCheck tracks the root domain (`wikipedia.org` , `doi.org` , `nih.gov`) and caps evidence items per domain, requiring that `STRONG_SUPPORT` be corroborated across at least **two distinct root domains**.

6. Relational Entailment & Contradiction Resolution

6.1 Relational Entailment vs. Lexical Overlap

HalluciCheck enforces strict semantic relation checking:

```
# Disambiguation of capital assertions vs financial homonyms
if "capital" in terms and "capital" in claim_lower:
    fin_capital = re.findall(
        r"\b(?:social|venture|working|human|physical)\s+capital\b|\bcapital\s+(?:expenditure|investment|gains|assets|cost|outlay)\b",
        evidence_text
    )
    all_capital = re.findall(r"\bcapital\b", evidence_text)
    if len(fin_capital) >= len(all_capital):
        # Exclude 'capital' from matched terms as it refers to finance, not sovereign seat of government
        matched = [t for t in matched if t != "capital"]
```

6.2 Authoritative Contradiction Weighting

When an assertion is contradicted:

- If a verified reference knowledge base or official registry (Wikipedia, Wikidata, CrossRef) directly documents that **Canberra** is the capital of Australia, or that Sydney is the **state capital of New South Wales**, the claim is flagged as `CONTRADICTION` (`X HALLUCINATED`).
- Spurious open-web text or unrelated mentions cannot dilute an authoritative contradiction into a vague "partial support" state.

7. Confidence Scoring & Clinical Risk Formulation

7.1 Quantitative Certainty Formulation

Individual claim certainty is calculated deterministically:

- ✓ VERIFIED**: Base confidence 88.0% – 92.0% with source diversity bonus +3.5% per additional independent root domain (capped at 98.0%).
- ⚠ SUSPICIOUS**: Base confidence 35.0% – 65.0% reflecting partial qualifier support or absence of peer-reviewed records.
- ✗ HALLUCINATED**: Base confidence 10.0% – 15.0% reflecting active contradiction by authoritative records.

7.2 Overall Session Confidence

Quantitative Session Score: Overall Score =
$$\frac{\sum_{i=1}^N W(S_i)}{N} \times 100$$

where:

Category	W(VERIFIED) = 1.0		W(SUSPICIOUS) = 0.45		W(UNVERIFIED) = 0.35	
Weights:	W(HALLUCINATED) = 0.0					

7.3 Risk Categorization Tiers

- HIGH CERTAINTY (≥ 80%)**: Low hallucination risk; all atomic propositions substantiated across independent domains.
- MODERATE RISK (50% – 79%)**: Secondary qualifiers lack verification or partial discrepancies observed.
- HIGH HALLUCINATION RISK (< 50%)**: One or more propositions contradicted by ground truth or total lack of substantiating evidence.

8. Frontend Verification Workspace & Human-in-the-Loop Diagnostics

The client workspace was designed in accordance with strict **Production Light Mode** standards:

1. Analysis Console:

- Paste-friendly `<textarea>` that never blocks pasting of multi-paragraph AI outputs.
- Real-time word and character counters.
- Quick-select sample preset chips.
- Keyboard shortcuts (`Ctrl+Enter` to run audit, `Ctrl+N` for new session).

2. Overall Certainty Ring & Metric KPIs:

- SVG certainty gauge using smooth `stroke-dashoffset` animation.
- Tabular numerals (`font-variant-numeric: tabular-nums`) preventing layout jitter.
- 4-card metric block: Total Claims, Verified Count, Suspicious Count, Hallucinated Count.

3. Interactive Highlighted Text Viewer:

- Renders the original input text sentence-by-sentence.
- Color-coded backgrounds with solid bottom border accents (Emerald for Verified, Amber for Suspicious, Rose for Hallucinated).
- Keyboard navigable (`tabIndex={0}` , `Enter` / `Space` to select, `Esc` to dismiss).

4. Deep Inspection Modal Card:

- Surfaces exact claim assertion, quantitative score, and 3-paragraph diagnostic rationale.
- Authority Registry Grid**: Displays verified domain tags and registry weights.
- Atomic Proposition Breakdown**: Itemizes every constituent proposition with support status.


- **Ground-Truth Citations:** Clickable direct URLs to source documents.

5. Color-Blind Accessibility (Triple Redundant Cues):

- No status relies solely on color. Every cue pairs color + icon (`CheckCircle2` , `AlertTriangle` , `XCircle`) + explicit uppercase label (`✓ VERIFIED` , `⚠ SUSPICIOUS` , `✗ HALLUCINATED`).

9. Cloud Deployment & DevOps Architecture

HalluciCheck is deployed with zero friction across multiple environments:



Production Deployments:

- **Live Vercel Production URL:** <https://aihallucicheck.vercel.app>
- **GitHub Repository:** <https://github.com/Niranchan67/aihallucicheck>
- **Unified Docker Server:** Multi-stage `Dockerfile` bundling Vite frontend + FastAPI backend on port `8000`.
- **Render Configuration:** Automated deployment manifest `render.yaml` ready for single-click hosting.

10. Empirical Benchmarks & Case Studies

Case Study A: Coordinate Sentence with False & True Claims

- **Input Text:**

"The capital of Australia is Sydney, and the human body consists of exactly 206 bones in an adult skeleton."

- **Verification Output:**

- **Proposition 1:** "The capital of Australia is Sydney."
→ `✗ HALLUCINATED` (14.0% Certainty)
Evidence Quote: "Sydney is the capital of the state of New South Wales and the largest city in Australia."
Discrepancy: Contradicted by official records documenting Canberra as the national capital.
- **Proposition 2:** "The human body consists of exactly 206 bones in an adult skeleton."
→ `✓ VERIFIED` (88.0% Certainty)
Evidence: Corroborated across biological and anatomical references.
- **Overall Score:** 50.0% (`MODERATE RISK`).
- **Client Rendering:** Sentence 1 highlighted in **Soft Red**; Sentence 2 highlighted in **Soft Green**.

Case Study B: Preserving Compound Nouns

- **Input Text:** "Water consists of hydrogen and oxygen."
- **Verification Output:** Kept as **1 single claim**; correctly substantiated with 92.0% certainty.

11. Security, Performance & Scalability

1. **Zero Auth Wall Policy:** The workspace provides immediate, friction-free verification without mandatory user login or paywalls.
 2. **Safe Database Operations:** Employs SQLAlchemy with connection recycling and auto-initializes tables in `/tmp/hallucicheck.db` to prevent read-only filesystem errors in serverless environments.
 3. **Network Resilience & Timeouts:** Outbound calls to external APIs use strict 5.0s timeouts with fallback mechanisms so that slow third-party servers never hang user audits.
 4. **Input Sanitization:** Automatically strips leading/trailing whitespace, truncates inputs exceeding character limits, and enforces parameter validation through Pydantic v2 schemas.
-

12. Conclusion & Future Outlook

HalluciCheck v2.0 represents a significant advancement in automated AI truth verification. By combining linguistic clause decomposition, multi-source academic retrieval, relational entailment, and transparent multi-paragraph reasoning, the system eliminates the blind spots of keyword matching and vector similarity.

Key Takeaways:

- Compound sentences are reliably dissected into independent assertions.
 - Contradictions are actively discovered and prioritized over spurious web matches.
 - End users receive comprehensive, transparent evidence proofs with direct clickable citations.
 - The platform is fully containerized and live in production on Vercel's global CDN.
-

Report compiled autonomously by Antigravity for HalluciCheck v2.0.
